

Q1. (C2 – 5 marks)

Explain the working of the **MapReduce model**. Identify its stages and advantages over traditional processing in distributed systems.

Working of MapReduce :

MapReduce works in two main steps:

1. **Map Step**
 - Breaks big data into smaller pieces
 - Each piece is processed separately (in parallel)
 - Outputs data as key-value pairs
2. **Reduce Step**
 - Gathers all key-value pairs with the same key
 - Combines or summarizes them into final results

Stages in MapReduce

1. **Splitting** – Input data is divided into chunks.
2. **Mapping** – Map function applied to each chunk.
3. **Shuffling** – Intermediate data grouped by key.
4. **Sorting** – Grouped keys are sorted for reduction.
5. **Reducing** – Final output is generated by Reduce function.
6. **Output** – Results written to storage (e.g., HDFS).

Advantages over Traditional Systems:

- Easy to write and understand
- Automatically handles parallelism and failures
- Scales to thousands of machines
- Efficient data processing using **data locality**

Q2. (C2 – 5 marks) Differentiate between **Distance Vector Routing (DVR)** and **Link State Routing (LSR)**. Explain how **Bellman-Ford** and **Dijkstra's algorithms** apply to these with basic **pseudocode**.

Distance Vector Routing (DVR)

- Routers share info **with neighbors only**.
- Uses **Bellman-Ford algorithm**.
- Routers know distance to other routers, **not full path**.
- Slower updates, can face issues like **count-to-infinity**.

🌐 Link State Routing (LSR)

- Routers share info **with all routers**.
- Uses **Dijkstra's algorithm**.
- Each router builds **full map** of the network.
- Faster and more accurate routing.

■ Bellman-Ford (used in DVR) – Finds shortest path by relaxing edges:

text

📄 Copy 🖋 Edit

```
Repeat (V - 1) times:
  For each edge (u, v):
    if distance[u] + weight < distance[v]:
      update distance[v]
```

■ Dijkstra's (used in LSR) – Finds shortest path using nearest unvisited node:

text

📄 Copy 🖋 Edit

```
Start from source:
  Pick node with smallest distance
  Update distances of its neighbors
  Repeat until all nodes are visited
```

Q3. (C2 – 5 marks) Compare shared memory and distributed memory architectures. Provide advantages, disadvantages, and suitable use-cases.

1. Shared Memory Architecture

Aspect	Description
Memory Access	All processors share the same physical memory
Communication	Through shared variables/memory
Speed	Fast communication

Aspect	Description
Complexity	Easier programming, harder synchronization

✓ *Advantages:*

- Simple to program and debug
- Faster data access and sharing

✗ *Disadvantages:*

- Hard to scale (limited number of processors)
- Risk of memory conflicts (need locks)

📌 *Use-cases:*

- Multicore processors
- Single-node parallel processing

🌐 2. Distributed Memory Architecture

Aspect	Description
Memory Access	Each processor has its own private memory
Communication	Via message passing (e.g., MPI)
Speed	Slower communication due to network delay
Complexity	More complex programming

✓ *Advantages:*

- Highly scalable
- No memory conflicts

✗ *Disadvantages:*

- Complex code (need explicit communication)

- Slower communication

✈ Use-cases:

- Supercomputers, cloud clusters
- Large-scale data processing (e.g., Hadoop, MPI)

Q4. (C3 – 5 marks) Explain how **barrier synchronization** works in MPI. Provide an example where synchronization is necessary.

◆ What is Barrier Synchronization?

In MPI, **barrier synchronization** means **all processes must stop and wait** at a certain point (the "barrier") until **every process reaches it**. After that, they all move forward **together**.

MPI provides this using:

```
MPI_Barrier(MPI_COMM_WORLD);
```

◆ Why It's Needed? (Example)

Imagine every process is loading data.

If one process starts using the data **before others finish loading**, it can cause errors.

So we use `MPI_Barrier()` to make sure:

- **All processes finish loading data**
- **Then start computing together**

◆ Simple Example:

```
// All processes load data
// Then wait at the barrier
MPI_Barrier(MPI_COMM_WORLD);
// All start computation together
```

Q5. (C2 – 5 marks) Name the **taxonomy** that categorizes computers into four types. Discuss the types briefly.

◆ **Name of the Taxonomy:**

Flynn's Taxonomy

◆ **Four Types of Computers in Flynn's Taxonomy:**

1. **SISD (Single Instruction, Single Data):**
 - One processor executes **one instruction** on **one data** at a time.
 - Example: Traditional single-core computer.
 2. **SIMD (Single Instruction, Multiple Data):**
 - One instruction is applied to **multiple data elements** at the same time.
 - Example: Graphics Processing Units (GPUs).
 3. **MISD (Multiple Instruction, Single Data):**
 - Many instructions operate on the **same data**.
 - Rare in practice; mostly used for **fault-tolerant systems**.
 4. **MIMD (Multiple Instruction, Multiple Data):**
 - Many processors execute **different instructions** on **different data**.
 - Example: Multi-core processors, parallel systems.
-

Q6. (C2 – 5 marks) Derive the formula for **scalability** that affects speedup in parallel processing and explain with a numeric example.

✓ **Step-by-Step Derivation:**

Let:

- **T_{serial}** = time taken by the program with 1 processor
- **P** = fraction of the program that can be parallelized
- **(1 – P)** = serial portion (cannot be parallelized)
- **N** = number of processors

1. Total Execution Time with Parallelism:

Saved memory full ⓘ

$$T_{\text{parallel}} = T_{\text{serial}} \times \left[(1 - P) + \frac{P}{N} \right]$$

2. Speedup Formula:

Speedup S is how much faster the program runs:

$$S = \frac{T_{\text{serial}}}{T_{\text{parallel}}}$$

Now substitute the parallel time from step 1:

$$S = \frac{T_{\text{serial}}}{T_{\text{serial}} \times \left[(1 - P) + \frac{P}{N} \right]}$$

Cancel out T_{serial} :

$$S = \frac{1}{(1 - P) + \frac{P}{N}}$$

🔴 Final Formula:

$$S = \frac{1}{(1 - P) + \frac{P}{N}}$$

This is Amdahl's Law — a key formula for analyzing scalability in parallel processing.

📘 Example:

Let:

- $P = 0.9$ (90% of the code is parallel)
- $N = 10$ processors

$$S = \frac{1}{(1 - 0.9) + \frac{0.9}{10}} = \frac{1}{0.1 + 0.09} = \frac{1}{0.19} \approx 5.26$$

Result:

With 10 processors and 90% parallel code, you get around **5.26× speedup**.
Even with more processors, the 10% serial part limits the maximum speedup.